
Field-agent modelling with CAMPO

Release 0.1

Computational Geography team UU

Dec 10, 2021

CONTENTS

1	Static Modelling	3
2	Spatio-temporal Modelling	11

Download this website as pdf.

Download this website as epub (for e-readers).

Software requirements

For the Campo labs you need to install [Campo](#) and have a GIS available for visualisation of results, preferably QGIS. If you install Campo following the instructions on the Campo website, QGIS is automatically installed in the same Conda environment as Campo, as it is in the list of dependencies.

Data sets

Course data sets can be downloaded:

- Static modelling, https://campo.computationalgeography.org/download/campo_static.zip
- Spatio-temporal modelling, https://campo.computationalgeography.org/download/campo_spatio_temporal.zip

Contents

STATIC MODELLING

1.1 Campo data and visualisation basics for point objects

1.1.1 Introduction & data set

In this exercise you will use the Campo Python module to execute static operations on fields and agents. You will use a data set of household locations and food outlets that also will be used in the spatio-temporal modelling exercises with Campo.

Campo is a model building framework for construction of field-agent based models, info is at <https://campo.computationalgeography.org>. You will have access to the capabilities of Campo through the Campo Python module. Under the hood, Campo uses several libraries. An important one is the LUE library for storing data containing both agents and fields. Information on LUE is at <https://lue.computationalgeography.org>. Note however that for the exercises below you will access LUE data always through Campo, so it is most important to understand Campo concepts.

Campo is still under development and in this exercise you will only use a subset of components of Campo as other components have not yet been documented and tested.

1.1.2 Phenomena, property sets and point objects

Open the template file `static_model.py` in your text editor and inspect it. The script imports the necessary Python modules and creates the static model class `FoodEnvironment`. You will add your model processes, i.e. operations on fields or agents, to the initial section of the model script. Processes stated in this section will be executed once when running the model.

We will model propensity for healthy food of households and use household locations and shop locations in the municipality of Utrecht, the Netherlands, for our model. The locations of stores and household are stored in two separate text files that we will use later on.

To add phenomena to a model you first need to initialize a model data object. Replace `pass` from the initial method and replace it by the following:

```
foodenv = campo.Campo()
```

Execute the script, it shouldn't print anything. The `foodenv` object will provide access to the methods for constructing agents, and eventually arrange the storage to the LUE dataset on disk.

A particular 'thing' that is represented by your model is represented by a 'phenomenon' in Campo. A phenomenon contains multiple agents or individuals, referred to as 'objects' in Campo. Examples of a phenomenon are trees, cars, catchments. Here we will create the phenomenon `foodstores`. To create a phenomenon, the `add_phenomenon` function is used. Add the phenomenon `foodstores` to the `foodenv` data object by typing the following statement in the initial part of the script:

```
foodstores = foodenv.add_phenomenon("foodstores")
```

Execute the script, again it shouldn't print anything.

Each object (agent, individual) in a phenomenon can have multiple properties attached to it, for instance color, weight, biomass. Properties can exist only at one location (x,y), for instance at the front door of a foodstore. These are so called point properties. Alternatively, properties can also exist in a spatial extent defined by a bounding box, which are called field properties. We will first focus on point properties. All point properties within a phenomenon are grouped together in a so called 'property set'.

Open the file `foodstores_frontdoor.csv` and have a look at its contents. Each row contains the (x,y) coordinates of a foodstore. Close the file without changing its content. We will create objects inside `foodstores` by reading from this file. Add the following statement:

```
foodstores.add_property_set("frontdoor", "foodstores_frontdoor.csv")
```

The first arguments assigns the name of the property set that is created and the second the input file name. Inside the `foodstores` phenomenon, Campo will create a number of objects equal to the number of lines in the input csv file, and a property set named `frontdoor` with a point location for each object (at the location of the coordinate given in the csv file).

1.1.3 Point properties

So far you have only defined the domain of the data set: the number of objects, and a property set (`frontdoor`) defining a location for each object. It is now time to add actual property values to this property set.

You can add properties by just using the dot notation on property sets and initialize them directly, e.g. setting the four digit `postal_code` to 1234 can be done by:

```
foodstores.frontdoor.postal_code = 1234
```

This will assign the value 1234 to the postal code property value of each object.

You can also assign random values drawn from a uniform distribution as follows:

```
foodstores.frontdoor.lower = -0.5
foodstores.frontdoor.upper = 0.5
foodstores.frontdoor.x_initial = campo.uniform(foodstores.frontdoor.lower, foodstores.
↪frontdoor.upper)
```

The first two statements set the lower and the upper bounds of the uniform distribution. The third statement draws a random value between these bounds for each object and assigns it to the property `x_initial`. This property refers to the quality of food (also, propensity) offered by the foodstore. Below zero implies unhealthy while above zero implies healthy food. You will in the spatio-temporal modelling exercise how this is used to model dietary habits over the city.

Question: What is a uniform distribution? Why would one use in certain cases a uniform distribution instead of the often used Gaussian distribution?

1.1.4 Writing data to disk and visualisation of point properties

All information in the LUE data set and properties you created so far exist only in memory. To store those to disk you need to perform a few additional steps at the end of your initial section. First create a LUE data set with the name `food_environment.lue`:

```
foodenv.create_dataset("food_environment.lue")
```

All phenomena, property sets and properties can then be stored in the LUE dataset by adding at the end of your initial section.

```
foodenv.write()
```

Run the script again and check if the `lue` file is stored on your hard disk. It should be in the same folder as your model.

You can open the `lue` file with the script `plot_point_objects.py`. Inspect it. The first function reads the data set from disk assigning it to `dataset`. The second function `campo.dataframe.select` can be used to read data from the `lue` dataset storing it in a standard Python data structure. Its first argument defines from which phenomenon data is read (here, `foodstores`), the second argument defines a list of strings defining which properties one would like to read, here `x_initial`. Add two lines at the bottom:

```
print(dataframe)
print(type(dataframe))
```

You will notice that for point objects, `dataframe` returned by the function is a dictionary showing the objects per row and property values per column.

For plotting you may want to export the contents of `dataframe` to a `.csv` file. You can do this by adding at the bottom:

```
campo.to_csv(dataframe, 'foodstores.csv')
```

Run the script and inspect the contents of the `csv` file. It contains the (x,y) coordinates of each object with the `x_initial` value added. Of course you can use the `csv` file to plot the data.

Alternatively you can export data to a GeoPackage file (GPKG) and then visualise in a GIS. To do so, add the following line to `plot_point_objects.py`:

```
campo.to_gpkg(dataframe, 'foodstores.gpkg', 'EPSG:28992')
```

It creates `foodstores.gpkg` which can for instance be displayed with QGIS. If you have QGIS installed in your Conda environment you can simply type:

```
qgis data/roads.gpkg foodstores.gpkg
```

It will open QGIS with the foodstore data and a roads layer for context. If QGIS is not in your Conda environment start it up like normal and then open the same files as layers by dragging them from the Browser window into the Layers window. In QGIS you can inspect property values with the Identify Features pointer (click on top of window on the button with the 'i' symbol). Or visualise them: right-click on 'foodstores' in the Layer window, select Properties, select Symbology, click on Symbol and change to Graduated, select for Value the property you want to show, for instance `x_initial`, click on Classify, click on OK. A filled colour is used to represent the value of the Property. Alternatively, you can select in the Graduated page also for Method: Size which shows the shops sized according to the Property value. Do not forget to click 'Classify'. For additional explanation refer to https://docs.qgis.org/2.18/en/docs/user_manual/working_with_vector/vector_properties.html.

1.2 Operations on point properties within a property set

1.2.1 Local operations

Local operations on point properties simply update the property values by using arithmetic operations, comparisons, etc. Open the script `static_point_operations.py` which is the script containing everything that you added in the exercises above. You will create now a propensity value `foodstores.frontdoor.x_value` by incrementing the initial value `x_initial` with a small random number. Add the following code to the script, just above the `create_dataset` function and run the script.

```
foodstores.frontdoor.lower_inc = 0.0
foodstores.frontdoor.upper_inc = 0.1
foodstores.frontdoor.increment = campo.uniform(foodstores.frontdoor.lower_inc,
↪ foodstores.frontdoor.upper_inc)

foodstores.frontdoor.x_value = foodstores.frontdoor.x_initial + foodstores.frontdoor.
↪ increment
foodstores.frontdoor.threshold = 0.0
foodstores.frontdoor.healthy = foodstores.frontdoor.x_value <= foodstores.frontdoor.
↪ threshold
```

Modify the `plot_point_objects.py` script to inspect the result and check the calculations.

1.3 Campo data and visualisation basics for field objects

1.3.1 Field properties and operations on field properties within a property set

In Campo, a phenomenon can include multiple property sets, where property sets are different regarding the spatial domain of the data. In the previous example you created a `frontdoor` property set inside `foodstores`. This had a point domain: each object was linked to a single location, the front door. Now we will add a second property set to the `foodstores` to simulate the spatial processes in the surrounding of each store. This is referred to as a field-agent.

Question: Give two other examples of the use of field-agents in spatial or spatio-temporal simulation.

Open the file `foodstores_surrounding.csv`. Each line gives the four corners of a bounding box surrounding a foodstore. The last two values give the number of rows and columns. In general, the spatial extents can be for all agents of the same size (e.g. fixed sized neighbourhoods), or the extents can differ between agents (e.g. different catchments). In our example we will consider the surrounding of the food stores with constant neighbourhood sizes per agent.

Open `static_model_fields.py`, it contains the main part used in the previous exercises.

In general, field agents are added comparably to the point agents with the `add_property_set` function. Add this function at the bottom of the script, reading `foodstores_surrounding.csv` and creating a new property set `surrounding`. Run the script to check if all is fine.

Now let's do the same `uniform` operation to create a field property inside the `surrounding` property set. Add this in the initial (at the bottom, but above the `write` statement):

```

foodstores.surrounding.lower = 3
foodstores.surrounding.upper = 12
foodstores.surrounding.c = campo.uniform(foodstores.surrounding.lower, foodstores.
↳ surrounding.upper)

```

Again run the script. In the next section you will inspect the result.

1.3.2 Writing data to disk and visualisation of field properties

To visualise field properties you can use `plot_field_objects.py`. Open the file and inspect it. Just like with the point properties, it opens the data set, and selects a property name, this time `c`. This is the property created in the previous section, a field property. The `to_tiff` function converts the field property value (here, `c`) for each object to a separate geotiff file, with file names `c_1.tiff`, `c_2.tiff`, etc. Run the script - it will print the content of `c`. Then, open some of the geotiff files (not all, as it may be too much for your computer) with a GIS, for instance QGIS. Again, if you have QGIS installed in your Conda environment, you could simply type the following command, which will open a subset of the files:

```
qgis data/roads.gpkg c_*0.tiff foodstores.gpkg
```

Just like with point objects, you can inspect cell values by selecting Identify Features and clicking on the map. Note that only values from an active geotiff file are plotted, to activate a layer, click on it (not uncheck/check as this determines whether the layer is displayed) in the Layers window. You may want to drag them all to a layer group which you can create in the Layers panel; this will allow you to check or uncheck them all (right-click on the layer group after dragging all geotiffs to the group). Refer to the QGIS manual for details https://docs.qgis.org/3.16/en/docs/user_manual/introduction/general_tools.html.

1.4 Operations on field properties within a property set

1.4.1 Distance calculation

With a two-dimensional raster-based spatial domain we particularly can execute spatial operations on the agents of a property set. You can, for example, calculate for each food store the distances to other foodstores within their spatial surrounding. This can be done with the `spread(start_locations, initial_friction, friction)` operation calculating for each cell the shortest distance to non zero cells.

First, you need to create a raster with locations of food stores in a surrounding. Use the `feature_to_raster` operation:

```

foodstores.surrounding.start_locations = campo.feature_to_raster(foodstores.surrounding,
↳ foodstores.frontdoor)

```

First input argument to the operation is the field property set holding the spatial surrounding of each food store, second argument are the front door locations of all the food stores in the modelling area. Return value of the operation is a raster with values of 1 for cells with corresponding food store locations and 0 otherwise. Note that this is a specific operation that has input arguments from two different property sets, `surrounding` and `frontdoor`.

Add the operation to the `static_model_fields.py` script. Run the script and visualise the output using `plot_field_objects.py` used in the previous section.

Now add the `spread` operation (syntax given above). You also need to create properties for the other two arguments of the `spread` operation, the initial friction distance (use the value 0) and the friction (use the value 1). Using these two values you calculate the absolute distances (in metres). Add the operation, run the script, and have a look at the result using `plot_field_objects.py`.

When you have this running, try creating a zone within 200 m from each food outlet

```
foodstores.surrounding.area = foodstores.surrounding.distance <= 200
```

Run the script and once more look at the result using `plot_field_objects.py`.

1.5 Operations between property sets: combining different phenomena and property sets

So far we mainly did calculations on property values within a certain property set of one and the same phenomenon. In many models however calculations are done between phenomena. An example is the `focal_agents` function which is similar to a focal, window, or convolution operation, but now applied on agents. The syntax is:

```
a.points.y = campo.focal_agents(a.points, a.fields.w, b.points.x)
```

It has inputs from two phenomena, `a` and `b`. The input `b.points.x` is a property belonging to the point property set `b.points`. For each object in `b`, `x` has a certain point value. The input `a.fields.w` is a property belonging to the field property set `a.fields`. For each object in `a`, property `b` is a raster where each raster cell contains a value. The input `a.points` merely sets the property set to return. The result is `a.points.y`, which is a property set in the point domain `a.points`. The function calculates for each object in `a`, the average value of all property values in the other phenomenon (`b.points.x`). The average value is a weighted average, where the weights are the values from `a.fields.w` at the location of each object in `b.points`. In an equation:

$$y_j = \frac{\sum_{i=1}^n w_j(x_i, y_i) x_i}{\sum_{i=1}^n w_j(x_i, y_i)}$$

The equation shows the calculation of the return value y_j , which is `a.points.y` for the j th object in `a`. It is used for each object j in `a`. It uses:

- x_i , the value of the i th object of `b.points.x`, and
- $w_j(x_i, y_i)$, the value of `a.fields.w` for the j th object in `a` at the location (x_i, y_i) of i th object of `b.points`.

To illustrate the use of `focal_agents`, consider the case where we are modelling the healthiness of food offered for sale at foodstores. The stores will adjust their product range to the dietary habits of their customers, that is, the households in the surroundings of a store. This mechanism can be modelled with the `focal_agents` operation. Open the script `static_two_phen.py`. It defines two phenomena, `foodstores`, and `households`. The households contain a property `households.frontdoor.x`, which is the households' propensity for eating healthy food. Here, we assign a random value between -2.0 and 2.0 to each household (we will expand upon this in the temporal modelling exercises!).

The foodstores contain two properties important for the calculation. The property `foodstores.surrounding.weight` gives for each store the weight in its surroundings (field-property). Be sure to understand how this weight is calculated (using the `spread` and `where` functions). The property `foodstores.frontdoor.y` is the healthiness of the food offered by the store, which is the average of the surrounding households, but weighted by `foodstores.surrounding.weight`.

Run the script and then inspect `plot_two_phen.py` which converts the outputs to files that you can view with QGIS, like in the previous exercises.

Open `foodstores.gpkg`, `households.gpkg`, as well as `area_143.tiff` and `weight_143.tiff` with QGIS. The two tiff files contain `foodstores.surrounding.area` and `foodstores.surrounding.weight` properties for object id 143 in the phenomenon `foodstores`. Check the calculation of `foodstores.frontdoor.y` by inspecting the y-value in `foodstores.gpkg` for object id 143. Note that it is a weighted sum of the x-values of `households.gpkg`, weighted by `weight_143.tiff`.

As you will have noticed the script assumes that a shop attracts customers up to a distance of 250 m away from the shop, this is defined in the operation calculating `foodstores.surrounding.area`. Change this value to 500 m and rerun the script. Again display the output for object id 143 and compare the results.

Question: What is approximately the range of values in the propensity of the foodstores? Compare this with the range in the propensity of the households. What is causing the observed difference?

Question: What is the effect of the size of the neighbourhood, for instance 250 or 500 m, on the resulting range in propensity values of the foodstores? Explain your answer.

SPATIO-TEMPORAL MODELLING

2.1 Point agents

In the spatio-temporal modelling exercises we will stepwise build a spatial model that simulates how the propensity for a healthy diet of households changes over time, in interaction with the healthiness of food offered at food stores. But we will first neglect the effect of the food stores. The propensity for healthy food of a household can be simulated with the following differential equation:

$$\frac{dx}{dt} = - \left(\frac{\beta_H}{1.0 + e^{-\gamma_H(x-a)}} - \frac{\beta_H}{2} \right)$$

In the equation, x is the propensity for healthy food of a household. The equation gives how it changes over time, as a function of the parameters β_H , γ_H , and a . The equation looks complicated but it is a simple curve. To plot the curve, open, inspect, and run `equations.py`. It creates `equations.pdf`. Neglect the second and third row of the pdf. The first plot in the first row is a plot of the equation above. The x-axis is x , the propensity of the household; the y-axis is dx/dt , the rate of change of the propensity over time. For propensity values with a positive rate of change, the propensity will increase over time, while with negative rates, it will decrease over time.

If you inspect the curve, you can see that the propensity will always tend to a value of zero in this case. This is a so called stable state. It is equal to the value of the parameter a , which we call the default propensity of the households. We assume for now it is the same for all households, zero. Propensities above zero refer to healthy food, while propensities below zero refer to unhealthy food. The idea here is that households will have a default value of the healthiness of their diet, which is set by the properties of the household itself (e.g. income, genes, etc). By external forcing, the propensity may change, but it will tend back to the default value. At least if we use only this equation!

Now it is time to run a spatio-temporal model that uses the equation to calculate how the propensity changes over time. Inspect `households.py`. It uses the same dataset from the static modelling exercises, but now defining all parameters from the equation above, in the initial section. In the dynamic section, then, the propensity, `self.hh.fdx` is updated for each timestep, using the equation given above. The time step duration is 3 months. We assume that at the start of the run, each household has a random propensity between -2 and 2.

Run the model. Inspect `plot_households.py`, which can be used to visualise the outputs. Just like in the static modelling exercises, it converts the output to gpkg files, but now for each timestep. They are written as `households_1.gpkg`, `households_2.gpkg`, etc. The second part of the script converts the propensity values x to a `.csv` file, which is then converted to a pdf. Each line in the pdf shows the change of the propensity over time of one household.

Run `plot_households.py` and open the `households_x.pdf` file. Also, open the gpkg files using QGIS. You could open them all as well as the road network from the folder data. In QGIS, visualise the values just like in the static modelling exercises. It is important to use the same Styles for each layer (time step) to enable comparison between timesteps. To do so, set the style for the first time step, `households_1`: right-click on the layer, select Properties, select Symbolology, click on Single symbol and change it to graduated. Then, select the Value to be plotted (x), use Method: color, click on Classify. Click on OK. The color of the symbols now represents the propensity. Now, you can copy the settings to the other layers: right-click on the `households_1` layer in the Layer panel and select Styles, Copy Style,

All Style Categories. Now you can paste the copied style to another layer, by right clicking on another layer. Do it for multiple layers. By selecting or deselecting layers you can show or hide them and inspect how values change over time.

Question: How does the value of the propensity change over time? What are the values at the last time step?

2.2 Field agents used for input in temporal model

Of course the dietary habits of households is not only influenced by internal household processes (represented in the previous section) but also by their influences from outside. For instance dietary habits in their social network or the healthiness of food offered in stores. If the food offered at stores is more healthy than a household's default propensity (represented by the parameter a) the actual propensity of the household will tend to be higher due to the influence of the food in the stores. This effect is modelled by our second model, which adds a second term, the effect of the healthiness of food offered at the store where the household is shopping. We refer to this food in the store as 'propensity' as well, represented by the symbol y in the equation below:

$$\frac{dx}{dt} = - \left(\frac{\beta_H}{1.0 + e^{-\gamma_H(x-a)}} - \frac{\beta_H}{2} \right) + \left(\frac{\beta_S}{1.0 + e^{-\gamma_S(y-a)}} - \frac{\beta_S}{2} \right)$$

To keep things simple, we use the same shape for the curve of the second term, however with a minus added and we will use slightly different parameter values. To inspect the curve of the second term, run `equations.py` and open the second pdf created by the script, `equation_food_store_effect.pdf`. Note that the x-axis gives propensity values of the food offered at a nearby store, the y-axis gives the corresponding rate of change in the propensity of the household (the outcome of the second term above). For now, we will assume that *all* stores offer food with a propensity of 0.2. As shown by the dashed orange line, this results in a rate of change (y-axis) of about 1.19 (year⁻¹). So the value of the second term for our example is 1.19.

To understand the values of the terms in the equation above, inspect `equations.pdf` again, in particular the second row. The first curve in the second row gives the value of the first term in the equation above, the so-called household effect. The second curve, then, gives the effect of the food outlet. For *all* household propensities (x-axis), the contribution to the rate of change is 1.19. The total rate of change, now, is the sum of the two curves, shown in the right plot (centre row). The position where the line crosses the x-axis is the equilibrium propensity of the households. As you can see, this is somewhat above 0.0, as a result of the influence of the food stores.

How does this influence the spatial pattern of household propensity? One would like to know how spatial patterns of propensity change driven by a certain spatial pattern of propensities of food stores. This can be evaluated with `households_foodstores.py`, defining from top to bottom:

- Two terms of the differential equation above.
- No changes have been made to the part of the script defining the starting values for the households.
- The foodstores phenomenon. As we do not have observed propensities available now, we generate food store propensities `self.fs.hd.y` that are either 0.0 (70% of them) or 0.2. The stores with 0.2 are randomly selected.
- The average of the propensity of food stores in the neighbourhood of each household, `self.hh.fd.y`. It is calculated using the focal operation just like we did in the static modelling exercises. The property `self.hh.fd.y` represents y in the equation above.
- The differential equation with the food store propensity effect added (dynamic section).

Run the script. Inspect `plot_households_foodstores.py` used for creating the output visualisations. Run the script and inspect the results.

Question: What is the highest value of the household propensity in the area at the end of the model run? Does it correspond with what you would expect from the value of the differential equation plot in `equations.pdf`, centre row, right plot?

2.3 Field agents in a temporal simulation

As a next step we will evaluate how households and food stores interact, leading to the emergence of spatial patterns in propensity. In the previous exercise we simulated the influence of the healthiness of food at foodstores on the healthiness of food consumed by households. We fixed the propensity of the foodstores. In reality, of course, food stores adjust the products offered to the products consumed. In this exercise we incorporate this mechanism by one simple addition to the model: for each time step the propensity of food stores is assumed to be equal to that of the propensity of the households in their surroundings. So, both food store propensity and household propensity are updated as a function of their surroundings, for each time step.

This concept leads to a full coupling of the households and foodstores. This can be understood with the help of the plots in the bottom row of `equations.pdf`. The first plot is still the same, it refers to the first term of the differential equation. As we assume now that food store propensity is equal to household propensity, the second plot changes. Instead of a fixed rate of change like in the previous exercise (centre row), the curve now follows the second term of the differential equation, where the rate of change changes with the propensity of the households (which is equal to that of the food stores in their surrounding). The total effect in the third column shows that two stable equilibria appear, one for propensities around -2 and another one for propensities around 2.

This approach has been coded in `households_foodstores_inter.py`. The main differences with the previous script are:

- Instead of fixing the food store propensity `self.fs.fd.y` it is now calculated as the average household propensity in the surroundings, in the dynamic section.
- As food store propensities change over time, `self.hh.fd.y` is now calculated for each time step as the average of the food store propensities in the surroundings of a household.

Run the model. It will take some time as the `focal_agents` operations now need to be done for each time step. Inspect `plot_households_foodstores_inter.py` to see what outputs are converted and run the script. View the output pdf files as well as the gpkg files with QGIS.

Question: What is the spatial pattern in the household propensities at the start, halfway, and at the end of the run? Explain the mechanisms that determine the spatial pattern and how it changes over time.
